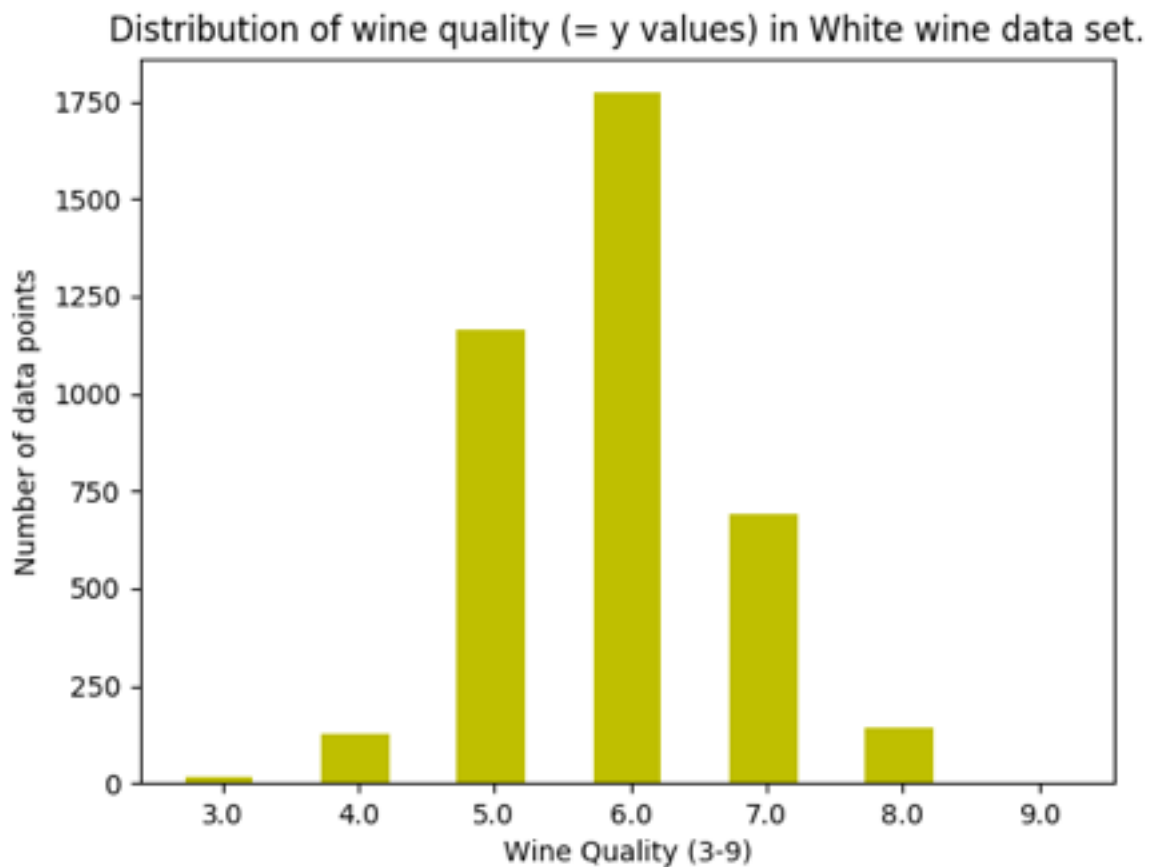


Machine Learning - Practical 1

Handin 1: Bar chart



Handin 2: MSE of Zero R

MSE of ZeroR on trainig set: 0.7768

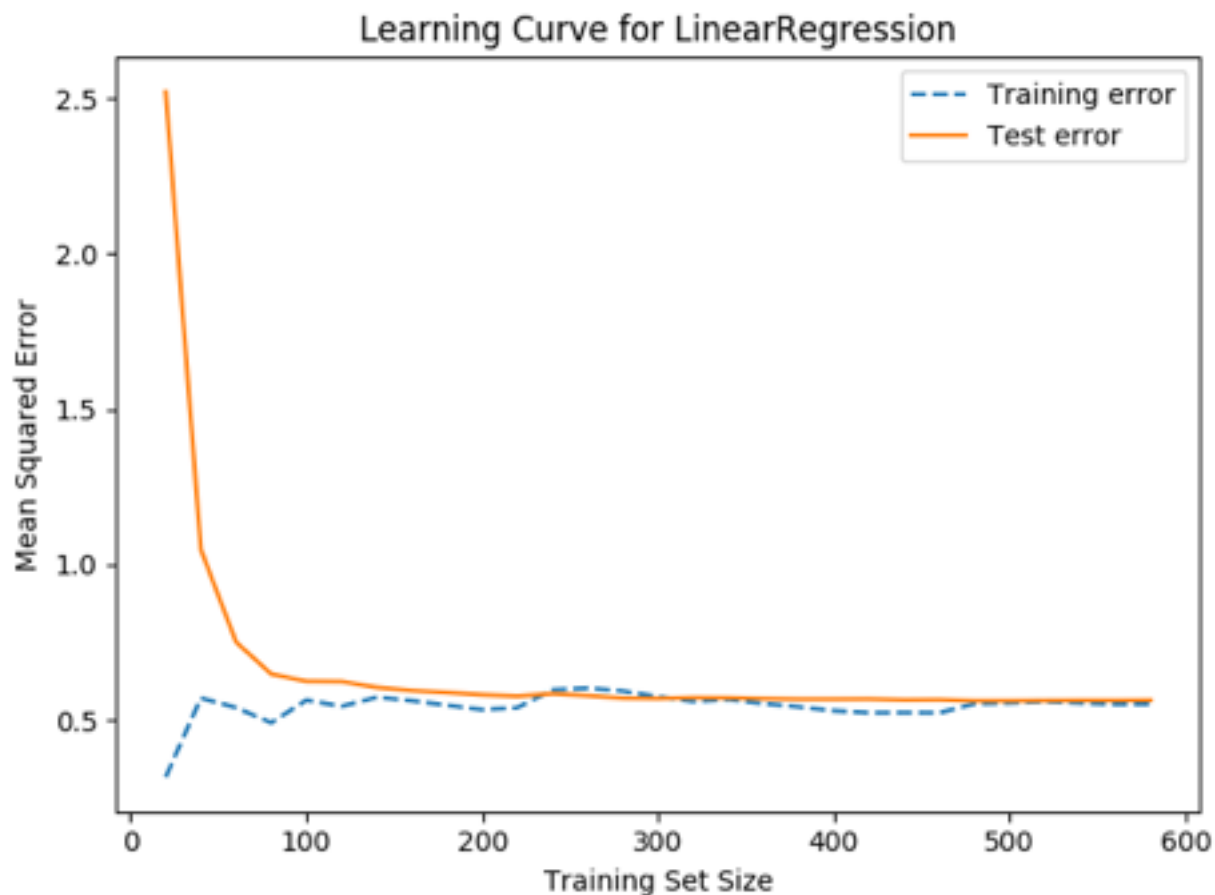
MSE of ZeroR on trainig set: 0.8139

Handin 3: MSE of Linear Model using Least Squares

MSE of LinReg on trainig set: 0.5640

MSE of LinReg on test set: 0.5607

Handin 4: Learning curves for Linear Model



Conclusion: The model soon stops overfitting (see increase of training error when size of training set grows from 20 to 40) and rapidly improves on the test error. It's hard to tell if the model is underfitting or perfectly fine, as we have no further knowledge of the underlying distribution (i.e., if it is perfectly linear, how much noise was added, ...) and therefore cannot judge just how well our model is fitting the data.